

# EECS E6892 Bayesian Models for Machine Learning

## Homework 1

John Min; jcm2199

February 27, 2014

### 1 Expectation Maximiation (E.M.) of Probabilistic PCA

$\{x_1, \dots, x_N\} \in \mathcal{R}^d, x_n \sim N(Wz_n, \sigma^2 I), W \in \mathcal{R}^{d \times k}$ .  $W$  is unknown and  $z_n \sim N(0, I)$ .

#### 1.1 Compute posterior of $z$

$$\begin{aligned}
 p(x_1, \dots, x_N | W^{\text{old}}) &\propto p(x_1, \dots, x_N | W^{\text{old}}, z_1, \dots, z_N) p(z_1, \dots, z_N) \\
 &= \prod_{n=1}^N \left[ \exp \left\{ \frac{(x_n - W^{\text{old}} z_n)^\top (x_n - W^{\text{old}} z_n)}{2\sigma^2} \right\} \exp \left\{ \frac{z_n^2}{2} \right\} \right] \\
 &\propto \exp \left\{ -\frac{1}{2} \sum_{n=1}^N \left[ z_n^\top z_n + \frac{1}{\sigma^2} z_n^\top W_{\text{old}}^\top W_{\text{old}} z_n - \frac{2}{\sigma^2} x_n^\top W_{\text{old}} z_n \right] \right\} \\
 &= \exp \left\{ -\frac{1}{2\sigma^2} \sum_{n=1}^N \left[ z_n^\top (\sigma^2 I + W_{\text{old}}^\top W_{\text{old}}) z_n - 2x_n^\top W_{\text{old}} z_n \right] \right\} \\
 &\sim \mathbf{N}(\mu, \Sigma) \\
 \mu &= \left( \frac{\sigma^2 I + W_{\text{old}}^\top W_{\text{old}}}{\sigma^2 I} \right)^{-1} W_{\text{old}}^\top X \\
 \Sigma &= \frac{\sigma^2 I}{\sigma^2 I + W_{\text{old}}^\top W_{\text{old}}}
 \end{aligned}$$

Note:

$$-\frac{1}{2}(x - \mu)^\top \Sigma^{-1}(x - \mu) = -\frac{1}{2}x^\top \Sigma^{-1}x + x^\top \Sigma^{-1}\mu + \text{const.}$$

Above,  $\mu = x^\top W$

$$z^\top \Sigma^{-1} \mu = z^\top W^\top x \Rightarrow \Sigma^{-1} \mu = W^\top x$$

#### 1.2 Take expectation of complete data log-likelihood

$$\begin{aligned}
 \mathbb{E}_{z_n} \left[ \ln p(x_1, \dots, x_N | W, z_1, \dots, z_N) \right] &= \sum_{n=1}^N \left[ \exp \left\{ \frac{x_n^2}{2} \right\} - 2\mathbb{E}[x_n^\top W z_n] + \mathbb{E}[z_n^\top W^\top W z_n] + \sigma^2 \mathbb{E}[z_n^\top z_n] \right] \\
 &= \sum_{n=1}^N \left[ \exp \left\{ \frac{x_n^2}{2} \right\} - 2x_n^\top W \mu + \mathbb{E}[(W z_n)^\top (W z_n)] + \sigma^2(\mu^2 + \Sigma) \right] \\
 &= \sum_{n=1}^N \left[ \exp \left\{ \frac{x_n^2}{2} \right\} - 2x_n^\top W \mu + [\text{Tr}(\mu^2 + \Sigma)]^\top [\text{Tr}(W W^\top)] + \sigma^2(\mu^2 + \Sigma) \right]
 \end{aligned}$$

#### 1.3 Maximize $Q(W, W^{\text{old}})$ w.r.t. $W$

Find the gradient with respect to  $W$  and set to 0:

$$\sum_{n=1}^N \left[ -2x_n^\top \mu + 2\text{Tr}(\mu^2 + \Sigma)W \right] = 0. \quad (\text{Note: } \nabla_A \text{Tr}(AB) = B^\top).$$

$$W^* \text{ aka } W^{\text{new}} = \left[ \text{Tr}(\mu^2 + \Sigma) \right]^{-1} \bar{X}^\top \mu \text{ where } \bar{X} = \frac{1}{N} \sum_{n=1}^N x_n$$

## 2 Implementation of Probabilistic Matrix Factorization

### 2.1 Maximum a posteriori (MAP)

Update steps:

$$u_i^{\text{MAP}} = (\lambda\sigma^2 I + V_i^\top V_i)^{-1} V_i^\top m_{u_i}$$

$$v_j^{\text{MAP}} = (\lambda\sigma^2 I + U_j U_j^\top)^{-1} U_j^\top m_{v_j}$$

Coordinate ascent is used for this MAP implementation where each update step for a particular  $u_i$  optimizes the joint log-likelihood for  $u_i$ , and the update step for  $v_j$  does the same for each  $v_j$ . Each update step, whether it be for  $u_i$  or for  $v_j$  is an  $\mathcal{L}_2$  regularized least squares solution, also known as ridge regression.

### 2.2 Gibbs sampling

$$\text{Sample: } u_i \sim N(\mu_i, \Sigma_i) \text{ where } \mu_i = (\lambda\sigma^2 I + V_i^\top V_i)^{-1} V_i^\top m_{u_i}, \Sigma_i = (\sigma^2 I + V_i^\top V_i)^{-1}$$

$$v_j \sim (N(\mu_j, \Sigma_j) \text{ where } \mu_j = (\lambda\sigma^2 I + U_j U_j^\top)^{-1} U_j^\top m_{v_j}, \Sigma_j = (\sigma^2 I + U_j U_j^\top)^{-1}$$

A burn-in period of 250 iterations is used after which we collect samples every 25 iterations. In each iteration, we sample each  $u_i$  from the posterior distribution with parameters that are being updated for each  $i$ . Then, we repeat for the  $v_j$ 's.

### 2.3 Similar Movies by $v_j$ using d=10

#### 2.3.1 Desperate Measures (1998)

Doors, The (1991); Mary Shelley's Frankenstein (1994); Malice (1993); Bye Bye, Love (1995); Love Affair (1994).

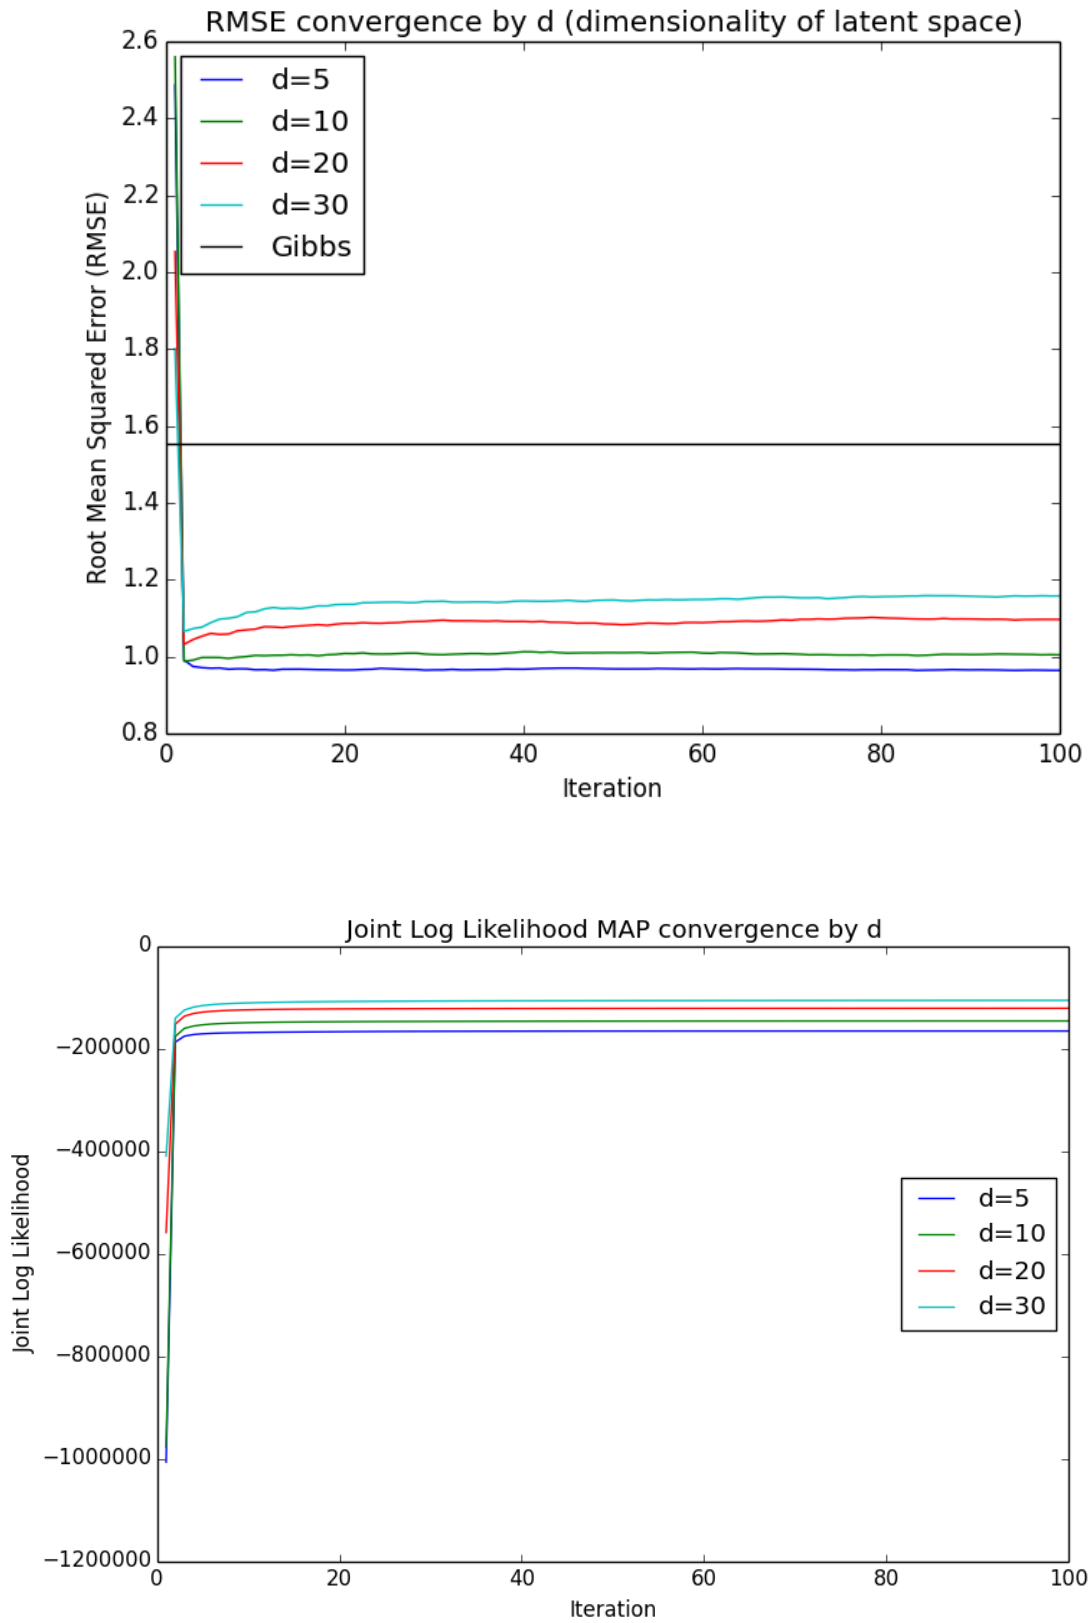
#### 2.3.2 Oliver & Company (1988)

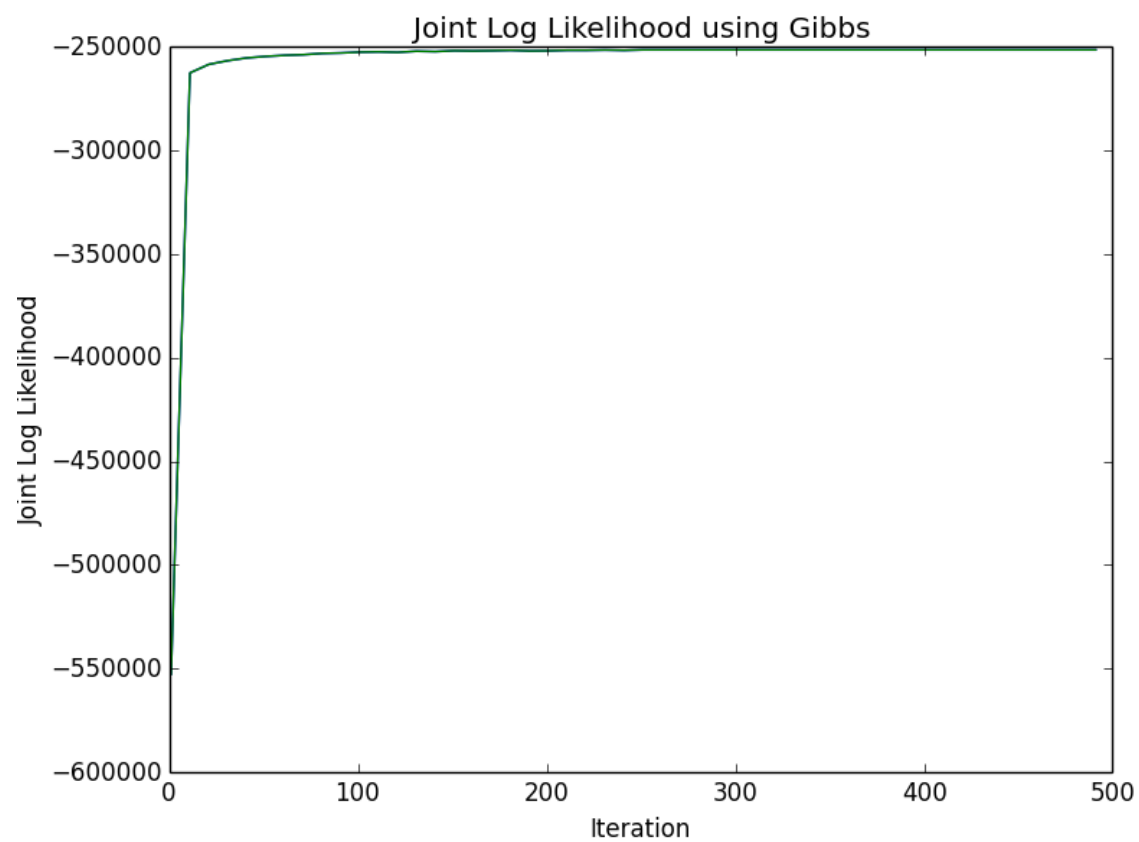
Indian in the Cupboard, The (1995); Prefontaine (1997); Abyss, The (1989); Kid in King Arthur's Court, A (1995); Air Up There, The (1994).

#### 2.3.3 Chinatown (1974)

Manhattan (1979); 8 1/2 (1963); Short Cuts (1993); Annie Hall (1977); Bonnie and Clyde (1967);

## 2.4 Plots of results





## 2.5 PMF - Python code

```
import numpy as np
import subprocess
import os
import matplotlib.pyplot as plt

def load_data_dict(filename):
    data = dict()
    with open(filename, 'rb') as f:
        for line in f:
            row = line.split(',')
            user, movie, rating = int(row[0])-1, int(row[1])-1, int(row[2])
            data[(user, movie)] = rating
    return data

def metadata(data, base_dir):
    users = list()
    movies = list()
    for user, movie in data.keys():
        users.append(user)
        movies.append(movie)
    users = sorted(list(set(users)))
    movies = sorted(list(set(movies)))
    N = len(users)
    M = int(subprocess.Popen('wc -l %s/movies.txt'%(base_dir),
        stdout=subprocess.PIPE, shell=True).communicate()[0].split()[0])
    #M = 1682
    return users, movies, N, M

def user_movie_dictionary(data, N, M):
    user_movie_dict = dict()
    for user_id in xrange(N):
        user_movie_dict[user_id] = list()
    for movie_id in xrange(M):
        try:
            data[(user_id, movie_id)]
            user_movie_dict[user_id].append(movie_id)
        except:
            pass
    return user_movie_dict

def movie_user_dictionary(data, N, M):
    movie_user_dict = dict()
    for movie_id in xrange(M):
        movie_user_dict[movie_id] = list()
    for user_id in xrange(N):
        try:
            data[(user_id, movie_id)]
            movie_user_dict[movie_id].append(user_id)
        except:
            pass
    return movie_user_dict

def dict_to_matrix(data_dict, N, M):
    mat = np.zeros((N, M))
    for user, movie in data_dict.keys():
        mat[user, movie] = data_dict[(user, movie)]
    return mat
```

```

#initialization
def initialize_factorization(N, M, d, sigma, lamb):
    U = np.zeros((N,d))
    V = np.zeros((d,M))
    I = np.identity(d)
    mean = np.zeros(d)
    cov = np.power(float(lamb),-1)*I
    for i in xrange(N):
        U[i,:] = np.random.multivariate_normal(mean, cov)
    for j in xrange(M):
        V[:,j] = np.random.multivariate_normal(mean, cov)
    return U, V, I

# user optimization
def update_user_MAP(M_train, U, V, N, I, user_movie_dict):
    for i in xrange(N):
        # compute V_i
        V_i = V[:, user_movie_dict[i]].transpose()
        # compute m_u
        m_u = M_train[i, user_movie_dict[i]]
        # compute u_MAP
        U[i,:] = np.dot(np.linalg.inv(lamb*np.power(float(sigma),2)*I +
            np.dot(V_i.transpose(), V_i)), np.dot(V_i.transpose(), m_u))
    return U

# movie optimizaiton
def update_movie_MAP(M_train, U, V, M, I, movie_user_dict):
    for j in xrange(M):
        # compute U
        U_j = U[movie_user_dict[j], :].transpose()
        # compute m_v
        m_v = M_train[movie_user_dict[j], j]
        # compute v_MAP
        V[:,j] = np.dot(np.linalg.inv(lamb*np.power(float(sigma),2)*I +
            np.dot(U_j, U_j.transpose()))), np.dot(U_j, m_v))
    return V

def predict_ratings(U, V):
    M_pred = np.dot(U,V)
    return M_pred

def compute_RMSE(M_pred, test):
    MSE_sum = 0
    N = len(test.keys())
    for i, j in test.keys():
        y_pred = round(M_pred[i,j])
        y = M_test[i,j]
        MSE_sum += np.power(y - y_pred, 2)
    RMSE = np.sqrt(MSE_sum/float(N))
    return RMSE

def log_likelihood(train, U, V, sigma, lamb, N, M):
    data_sum = 0
    u_sum = 0
    v_sum = 0
    for i, j in train.keys():
        data_sum += np.power(train[i,j] - np.dot(U[i,:].transpose(), V[:,j]), 2)
    for i in xrange(N):
        u_sum += np.dot(U[i,:].transpose(), U[i,:])
    for j in xrange(M):

```

```

v_sum += np.dot(V[:,j].transpose(), V[:,j])
joint_LL = np.power(sigma,-2)*-.5*data_sum - .5*lamb*u_sum - .5*lamb*v_sum
return joint_LL

# coordinate ascent
def coordinate_ascent(M_train, test, N, M, d, N_iterations, user_movie_dict, movie_user_dict, sigma, lamb):
    U, V, I = initialize_factorization(N, M, d, sigma, lamb)
    rmse_list = list()
    log_likelihood_list = list()
    for iter in xrange(N_iterations):
        U = update_user_MAP(M_train, U, V, N, I, user_movie_dict)
        V = update_movie_MAP(M_train, U, V, M, I, movie_user_dict)
        M_pred = predict_ratings(U, V)
        rmse = compute_RMSE(M_pred, test)
        rmse_list.append(rmse)
        joint_LL = log_likelihood(train, U, V, sigma, lamb, N, M)
        log_likelihood_list.append(joint_LL)
    return U, V, rmse_list, log_likelihood_list

def plot_RMSE(N_iterations, rmse):
    X = [x+1 for x in xrange(N_iterations)]
    plt.plot(X, rmse)
    plt.xlabel('Iteration')
    plt.ylabel('RMSE')
    plt.show()

## GIBBS SAMPLING
def update_user_Gibbs(M_train, U, V, N, I, user_movie_dict):
    for i in xrange(N):
        # compute V_i
        V_i = V[:, user_movie_dict[i]].transpose()
        # compute m_u
        m_u = M_train[i, user_movie_dict[i]]
        # compute u_MAP
        mean_user = np.dot(np.linalg.inv(lamb*np.power(sigma,2)*I +
            np.dot(V_i.transpose(), V_i)), np.dot(V_i.transpose(), m_u))
        cov_user = np.linalg.inv(lamb*I + np.power(sigma,-2)*np.dot(V_i.transpose(), V_i))
        U[i,:] = np.random.multivariate_normal(mean_user, cov_user)
    return U

def update_movie_Gibbs(M_train, U, V, M, I, movie_user_dict):
    for j in xrange(M):
        # compute U_j
        U_j = U[movie_user_dict[j], :].transpose()
        # compute m_v
        m_v = M_train[movie_user_dict[j], j]
        # compute v_MAP
        mean_movie = np.dot(np.linalg.inv(lamb*np.power(sigma,2)*I +
            np.dot(U_j, U_j.transpose()))), np.dot(U_j, m_v))
        cov_movie = np.linalg.inv(lamb*I + np.power(sigma,-2)*np.dot(U_j, U_j.transpose()))
        V[:,j] = np.random.multivariate_normal(mean_movie, cov_movie)
    return V

def initialize_gibbs_dict(test):
    gibbs_dict = dict()
    for i,j in test.keys():
        gibbs_dict[(i,j)] = list()
    return gibbs_dict

def sample_gibbs(train, gibbs_dict, U, V):

```

```

    for i, j in train.keys():
        gibbs_dict[i,j].append(np.dot(U[i,:],V[:,j]))
    return gibbs_dict

def compute_RMSE_Gibbs(gibbs_dict, test):
    MSE_sum = 0
    N = len(test.keys())
    for i, j in test.keys():
        y_pred = round(np.mean(gibbs_dict[i,j]))
        y = M_test[i,j]
        MSE_sum += np.power(y - y_pred, 2)
    RMSE = np.sqrt(MSE_sum/float(N))
    return RMSE

def gibbs(M_train, train, test, N, M, d, sigma, lamb, N_gibbs, burn_in, thinning):
    U, V, I = initialize_factorization(N, M, d, sigma, lamb)
    gibbs_dict = initialize_gibbs_dict(test)
    log_likelihood_list = list()
    for iter in xrange(burn_in):
        U = update_user_Gibbs(M_train, U, V, N, I, user_movie_dict)
        V = update_movie_Gibbs(M_train, U, V, N, I, movie_user_dict)
    if iter%10 == 0:
        joint_LL = log_likelihood(train, U, V, sigma, lamb, N, M)
        log_likelihood_list.append(joint_LL)
    for iter in xrange(N_gibbs - burn_in):
        if iter%10 == 0:
            joint_LL = log_likelihood(train, U, V, sigma, lamb, N, M)
            log_likelihood_list.append(joint_LL)
        if iter%thinning == 0:
            gibbs_dict = sample_gibbs(test, gibbs_dict, U, V)
            rmse = compute_RMSE_Gibbs(gibbs_dict, test)
    return rmse, log_likelihood_list

if __name__ == "__main__":

    # parameterization
    sigma = np.sqrt(0.25)
    lamb = 10
    d_list = [10, 20, 30]
    d = 10

    base_dir = os.path.join(os.getcwd(), 'movie_ratings')
    train = load_data_dict(os.path.join(base_dir, 'ratings.txt'))
    test = load_data_dict(os.path.join(base_dir, 'ratings_test.txt'))
    users, movies, N, M = metadata(train, base_dir)

    user_movie_dict = user_movie_dictionary(train, N, M)
    movie_user_dict = movie_user_dictionary(train, N, M)

    M_train = dict_to_matrix(train, N, M)
    M_test = dict_to_matrix(test, N, M)

    N_iterations = 100

    N_gibbs = 500
    burn_in = 250
    thinning = 25

```



```

d=5
U5, V5, rmse_5, LL_5 = coordinate_ascent(
    M_train, test, N, M, d, N_iterations, user_movie_dict, movie_user_dict, sigma, lamb)
d=10
U10, V10, rmse_10, LL_10 = coordinate_ascent(
    M_train, test, N, M, d, N_iterations, user_movie_dict, movie_user_dict, sigma, lamb)
d=20
U20, V20, rmse_20, LL_20 = coordinate_ascent(
    M_train, test, N, M, d, N_iterations, user_movie_dict, movie_user_dict, sigma, lamb)
d=30
U30, V30, rmse_30, LL_30 = coordinate_ascent(
    M_train, test, N, M, d, N_iterations, user_movie_dict, movie_user_dict, sigma, lamb)
rmse_gibbs, LL_gibbs = gibbs(
    M_train, train, test, N, M, d, sigma, lamb, N_gibbs, burn_in, thinning)

X_coord = [x+1 for x in xrange(N_iterations)]
X_gibbs = [10*x+1 for x in xrange(len(LL_gibbs))]

plt.plot(X_coord, rmse_5, label='d=5')
plt.plot(X_coord, rmse_10, label='d=10')
plt.plot(X_coord, rmse_20, label='d=20')
plt.plot(X_coord, rmse_30, label='d=30')
plt.axhline(y=rmse_gibbs, color='black', label='Gibbs')
plt.xlabel('Iteration')
plt.ylabel('Root_Mean_Squared_Error_(RMSE)')
plt.title('RMSE_convergence_by_d_(dimensionality_of_latent_space)')
plt.legend(loc=2)
plt.savefig('RMSE.png')
plt.show()

plt.plot(X_coord, LL_5, label='d=5')
plt.plot(X_coord, LL_10, label='d=10')
plt.plot(X_coord, LL_20, label='d=20')
plt.plot(X_coord, LL_30, label='d=30')
plt.xlabel('Iteration')
plt.ylabel('Joint_Log_Likelihood')
plt.title('Joint_Log_Likelihood_MAP_convergence_by_d')
plt.legend(loc=5)
plt.savefig('JLL.png')
plt.show()

plt.plot(X_gibbs, LL_gibbs)
plt.xlabel('Iteration')
plt.ylabel('Joint_Log_Likelihood')
plt.title('Joint_Log_Likelihood_using_Gibbs')
plt.savefig('JLL_Gibbs.png')
plt.show()

U = U10
V = V10

f = open(base_dir+'/movies.txt')
i = 0
movies = dict()
for line in f:
    movie_name = line.split('\n')[0]
    movies[i] = movie_name
    i += 1

import pandas as pd

```

```

import random

n = range(M)
random.shuffle(n)

movie_list = n[:3]

sim_movie_dict = dict()
for m in movie_list:
    sim_movie_dict[m] = list()
    print 'Base_Movie:_', movies[m]
    col = range(M)
    col.remove(m)
    dist_Euc = []
    for j in col:
        d = np.linalg.norm(V[:,j] - V[:,m])
        dist_Euc.append(d)
    dist_Euc = pd.Series(dist_Euc)
    dist_Euc.sort() #inplace sort
    for idx in dist_Euc.index[:5]:
        sim_movie_dict[m].append(col[idx])
    print '5_most_similar_movies:'
    for x in sim_movie_dict[m]:
        print movies[x]

```